

## 4 RESULTS

### 4.1 Test Regime

The "production\_campaign" featured two main pushes of testing. The first in December 2015 through to January 2016 tested Google Maps Engine endpoints before their shutdown. The campaign's main thrust took place in March 2016 where the majority of tests queried the old OGC endpoints and the new Esri endpoints.

Table 1 - Subtest counts

The user initiated 284 TestMasters resulting in 16,144 TestEndpoints with similar numbers of LocationTests, NetworkTests and PingTests.

There were three theatres of action in the campaign. Each test is mapped in a

Leaflet web map using the location of its Vector's PreTestLocation (the LocationTest completed before the TestEndpoint began). Visualising 16,000 points would result in an ineffective map, so here closely clustered points are generalised into a heat map. A beneficial side effect of generalisation is to obfuscate precise locations.

The majority of tests took place in Sydney, NSW and its environs. In particular the regular commute over the harbour to the Central Business District, and the roads and freeways to neighbouring cities.

The figures look great. You need to also explain the figures within the text. The reason for this is to ensure that the reader understands from you what you want them to understand.

Also when you explain the figures in the text it is important that you refer to them by name.

The above comments also are for tables.

With the charts, always speak =& refer to the ones that are in the paper, not to the latest ones online.

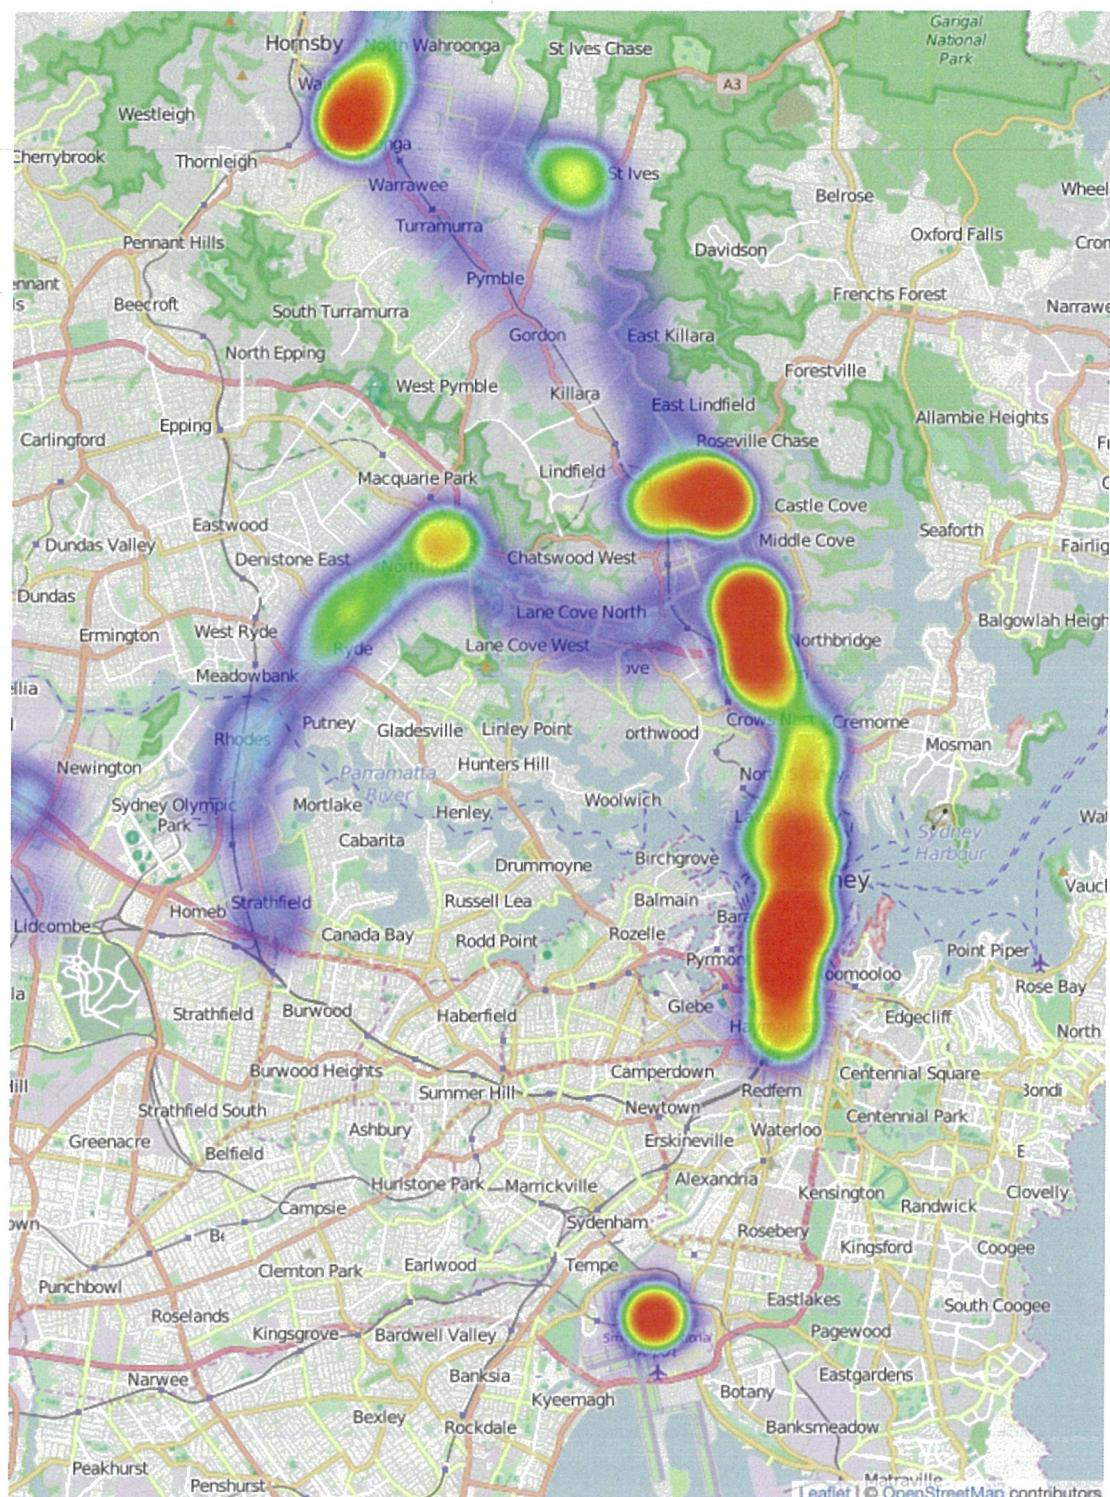


Figure 2 - Sydney Test Heat Map, basemap tiles copyright OpenStreetMap Contributors

Several discrete bursts of tests took place in Bathurst, NSW and the highway back and forth to Sydney, NSW.

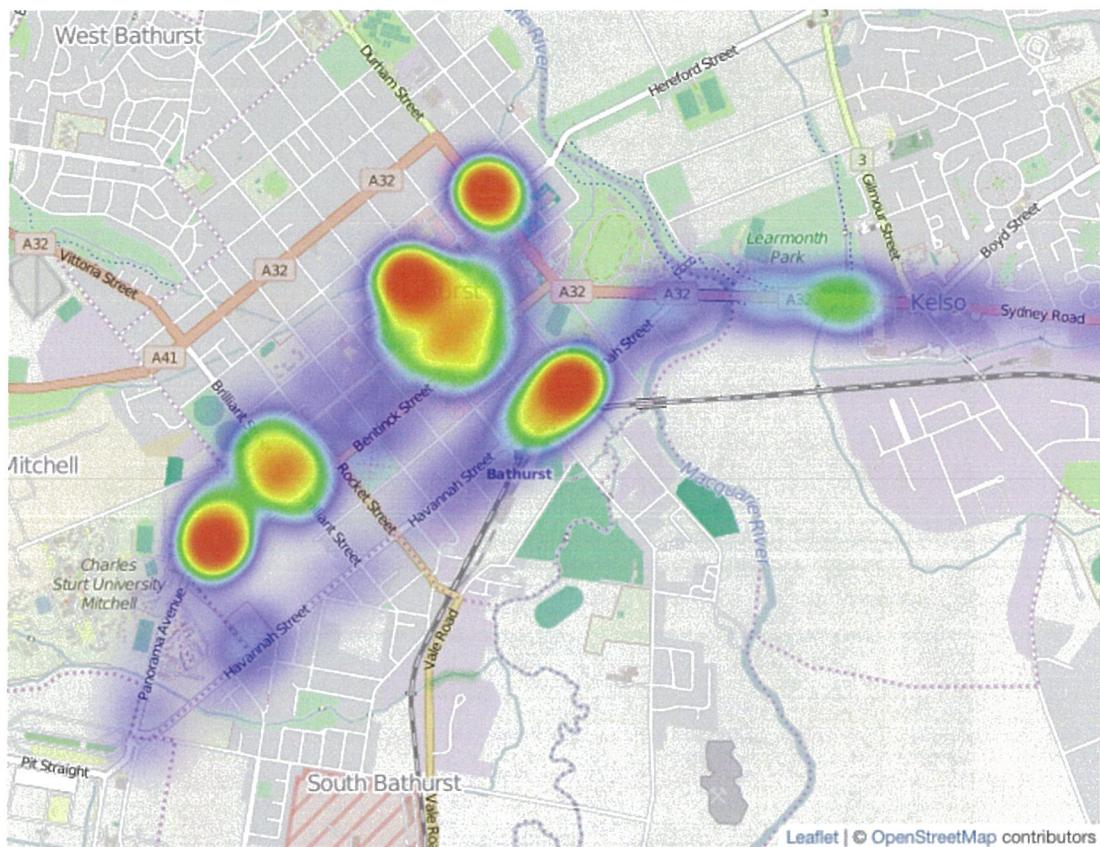


Figure 3 - Bathurst Test Heat Map, basemap tiles copyright OpenStreetMap Contributors

Townsville, QLD was the theatre with the least number of tests, but some interesting mobile situations involving ferry crossings and steep terrain on Magnetic Island.

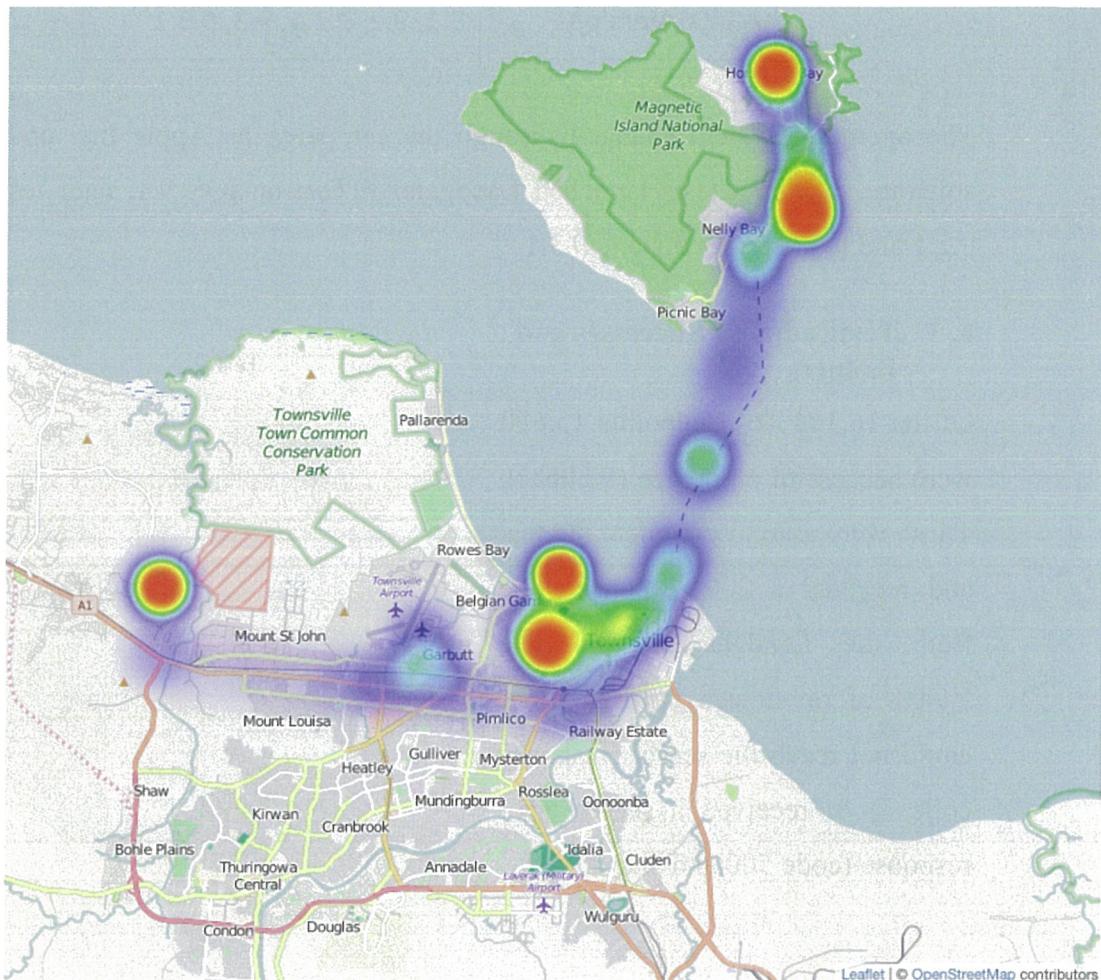


Figure 4 - Townsville Test Heat Map, basemap tiles copyright OpenStreetMap Contributors

#### 4.2 Test Device Hardware

All tests were performed on an Apple iPhone 6S, model A1688 (a.k.a. iPhone8,1), with 64GB of storage. The standard device comes with a range of mobile radios across many bands; LTE, HSDPA, CDMA, GSM, EDGE, Wi-Fi radios a/b/g/n/ac and GPS and GLONASS receivers {Anonymous:uf}.

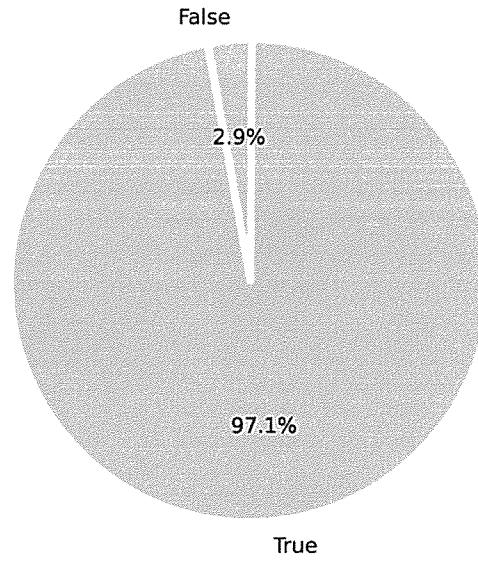
Table 2 - Test device characteristics captured by LandgateAPITest iOS application

<b>Campaign Name</b>	production_campaign
<b>All Device Types</b>	iPhone8,1
<b>All iOS Versions</b>	9.1, 9.2, 9.2.1, 9.3, 9.3.1

The operating system changed through the campaign as Apple Inc. updated their software. The first tests launched LandgateAPITest on iOS 9.1, later tests on 9.2, 9.2.1 and later still on 9.3 and 9.3.1.

#### 4.3 TestEndpoint Successes and Failures

Of the 16,144 TestEndpoints 15,670 were successful on device (97.06%). These were able to complete the test and received a 200 response code from the Landgate server. The 2.94% of on device failures either could not reach the server (response code 0) or received a server error response (code 500 and above).



LandgateAPITest's /Analyse function compared each

Figure 5 - TestEndpoints successful and failed on device

TestEndpoint's response data to the stored reference data and determined that 13,220 of them match, setting the resultant Vector's referenceCheckSuccess flag to True.

Closer examination of referenceCheckSuccess by test type showed 9 test types that consistently failed their reference checks (less than 5% passed). All such Vectors had their ReferenceCheckValid flag set to False to exclude them en masse from further analysis on the assumption that there was a systematic issue with their reference data.

Notably, the GetCapabilities tests rarely passed reference checks. Likely causes include changes to services offered during the test campaign or possibly timestamps buried in the XML response and reference conflicting.

Table 3 - Percentage of reference checks successful

Test Name	Percent Reference Check Successful
ESRI - BusStops - AttributeFilter - GET - JSON	98.79%
ESRI - BusStops - AttributeFilter - POST - JSON	98.20%
ESRI - BusStops - Big - GET - JSON	101.06%
ESRI - BusStops - Big - POST - JSON	101.03%
ESRI - BusStops - FeatureByID - GET - JSON	99.00%
ESRI - BusStops - FeatureByID - POST - JSON	98.76%
ESRI - BusStops - GetCapabilities - GET - JSON	97.64%
ESRI - BusStops - GetCapabilities - POST - JSON	98.99%
ESRI - BusStops - IntersectFilter - GET - JSON	97.31%
ESRI - BusStops - IntersectFilter - POST - JSON	98.21%
ESRI - BusStops - Small - GET - JSON	96.98%
ESRI - BusStops - Small - POST - JSON	97.89%
ESRI - Topo - Big - POST - Image	99.45%
ESRI - Topo - Small - GET - Image	99.32%
ESRI - Topo - Small - POST - Image	100%
GME - AerialPhoto - Big - GET - Image	96.97%
GME - AerialPhoto - GetTileKVP - GET - Image	95.65%
GME - AerialPhoto - GetTileKVP2 - GET - Image	93.75%
GME - AerialPhoto - GetTileKVP3 - GET - Image	90.63%
GME - AerialPhoto - GetTileKVP4 - GET - Image	90.91%
GME - AerialPhoto - Small - GET - Image	86.36%
GME - AerialPhoto - WMSGetCapabilities - GET - XML	0%
GME - AerialPhoto - WMTSGetCapabilities - GET - XML	0%
GME - BusStops - AttributeFilter - GET - JSON	83.87%
GME - BusStops - Big - GET - JSON	0%
GME - BusStops - DistanceFilter - GET - JSON	93.75%
GME - BusStops - FeatureByID - GET - JSON	96.55%
GME - BusStops - IntersectFilter - GET - JSON	90.63%
GME - BusStops - Small - GET - JSON	0%
OGC - AerialPhoto - GetTileKVP - GET - Image	98.85%
OGC - AerialPhoto - GetTileRestful - GET - Image	98.36%
OGC - BusStops - AttributeFilter - GET - JSON	0.83%
OGC - BusStops - AttributeFilter - GET - XML	99.72%
OGC - BusStops - AttributeFilter - POST - JSON	98.63%
OGC - BusStops - AttributeFilter - POST - XML	99.45%
OGC - BusStops - Big - GET - JSON	100%
OGC - BusStops - Big - GET - XML	100%
OGC - BusStops - Big - POST - JSON	100%
OGC - BusStops - Big - POST - XML	100%

??

OGC - BusStops - FeatureByID - GET - JSON	99.30%
OGC - BusStops - FeatureByID - GET - XML	98.85%
OGC - BusStops - FeatureByID - POST - JSON	100%
OGC - BusStops - FeatureByID - POST - XML	99.54%
OGC - BusStops - GetCapabilities - GET - XML	2.88%
OGC - BusStops - GetCapabilities - POST - XML	1.26%
OGC - BusStops - IntersectFilter - GET - JSON	99.73%
OGC - BusStops - IntersectFilter - GET - XML	99.46%
OGC - BusStops - IntersectFilter - POST - JSON	99.46%
OGC - BusStops - IntersectFilter - POST - XML	100%
OGC - BusStops - Small - GET - JSON	98.81%
OGC - BusStops - Small - GET - XML	100%
OGC - BusStops - Small - POST - JSON	99.76%
OGC - BusStops - Small - POST - XML	98.83%
OGC - Topo - Big - GET - Image	3.63%
OGC - Topo - Small - GET - Image	3.48%

See Appendices A and B for a list of URLs for each request.

Of the remaining tests only 79 (0.6%) failed a reference check. And 2.9% of the referenceCheckValid tests failed on the device itself.

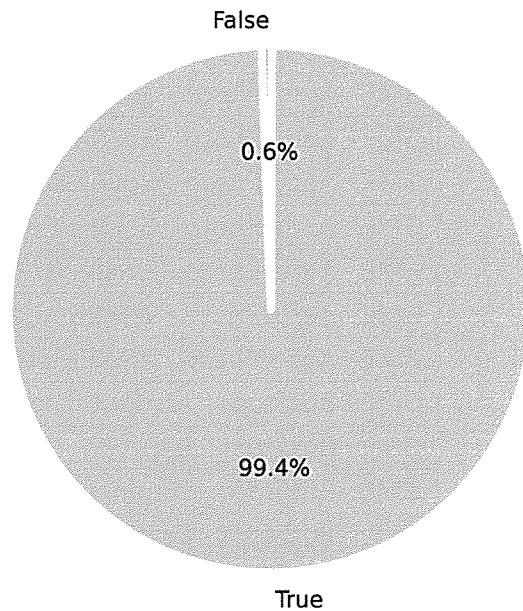


Figure 6 - Percentage of Vectors with failed reference checks

#### 4.4 Test Results by Response Time

The LandgateAPITest web application produces charts of current data on request to the /graph endpoint. All charts in this text are saved versions of these as they stood on the 7th of May, 2016. Links are provided in the text to the endpoint for each chart displayed so that the reader may receive the latest information.

The various requests are subcategorised by their test name, a general description denoting near identical requests across the three server types. A FeatureByID request returns the same data from all three servers, though it may not be in the same format (GML, Esri JSON, GeoJSON for example). This pie chart was modified to exclude some of the smallest percentage test types to aid reading clarity.

Figure 7...

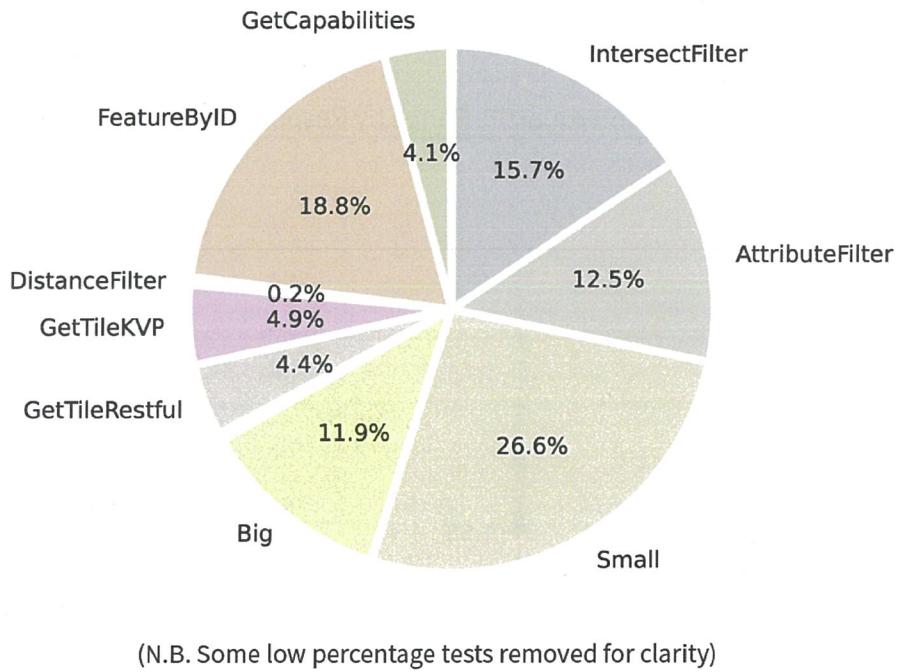


Figure 7 - Percentage of Vectors by test type N.B. smaller categories removed for clarity

Aligned with Fowler, Hameseder and Peterson's (2012) experimental control showing that response data size affects response time, LandgateAPITest requests "Small" and "Big" responses. Small requests are either for a few features in GML or JSON or an image only a few tens of pixels in dimension. Big requests ask for 100 vector features or images 500 pixels in dimension.

The distribution of their response times are shown in a box and whiskers chart, available at this [link](#). Box and whiskers charts show the interquartile range of a distribution. The dataset is divided into four equal parts around the median value, shown as a red line. The first quartile (Q1) to the third quartile (Q3) are contained within the blue box and contain 50% of the points in the dataset. The "whiskers"

actually have the link there

above and below the box show the range of the data and in a normal distribution would contain over 99% of the data points. Skewed distributions end up excluding the outliers from the interquartile range, shown here as blue crosses.

The "Big" requests have a similar Q1 to Q3 (interquartile range) to "Small" ones. The lowest values in the whiskers are significantly slower to arrive. Both have a significant number of outliers above the maximum response time whisker.

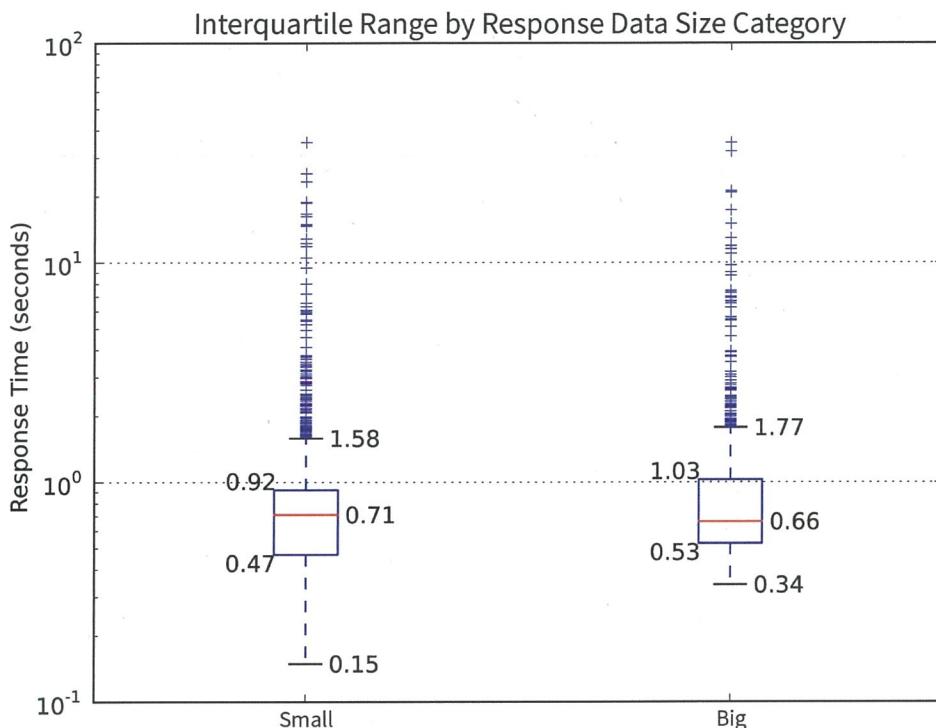


Figure 8 - A subset of test types where the request calls for either a small sized response or a larger one comparing their response times

Spatial servers can filter results either by a function of the attributes of each feature, returning features from any location meeting certain criteria of their properties. Features may instead be filtered by a spatial function, returning features from a specific location of any attribute value. The response time frequency distributions for four test types which call upon the server to filter results are shown in the box and whiskers chart. Feature by ID calls for a single feature with an exactly matching ID, an Attribute Filter test requests features with a text "location" property containing the

word "Curtin". Spatial intersect requests provide an envelope (minX, min Y, max X and max Y) covering the Curtin University Bentley campus and request only features intersecting the envelope. The Distance Filter was only requested from GME servers, returning only the closest feature to a point within Curtin University's Bentley campus.

The two attribute filters generally show a distribution of response times shorter than the two spatial filters. The confidence in this result is not great. Firstly, all have a significant number of high outliers denoting skewed distributions. The spatial filter medians are only 2 to 3 tenths of a second slower than the two attribute filters. The much smaller Distance Filter sample size makes it less worthy of consideration.

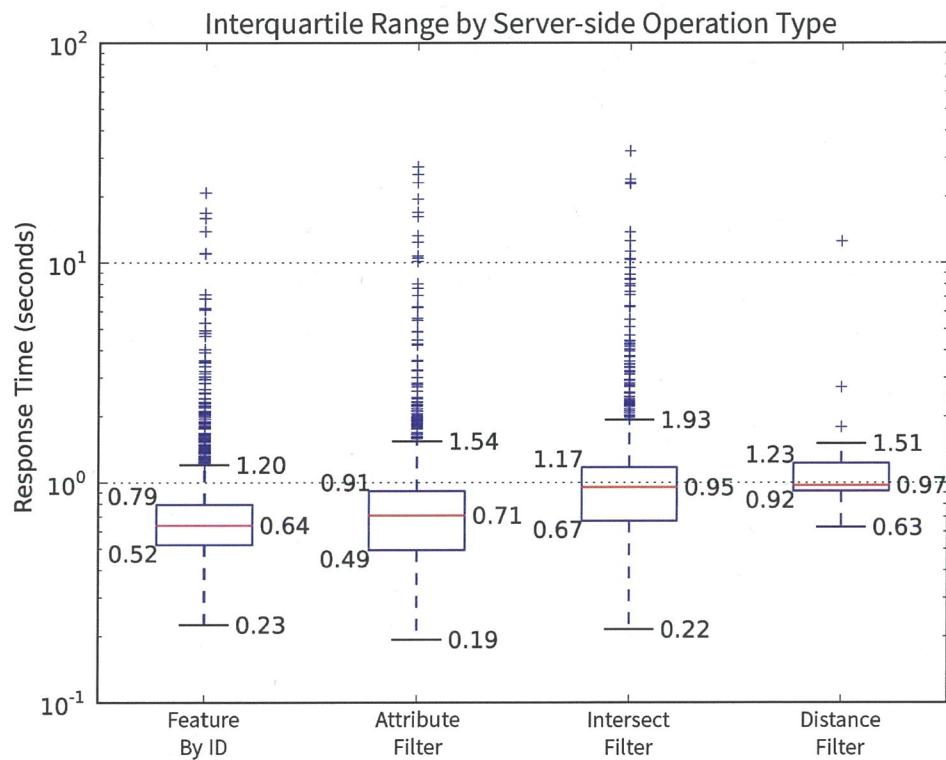
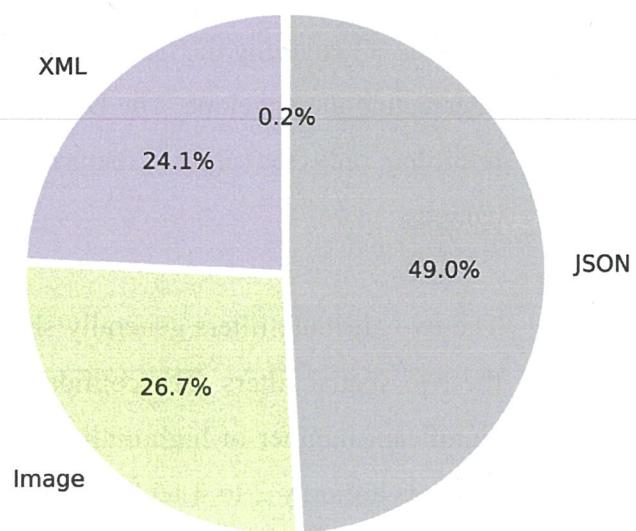


Figure 9 - A subset of test types which call upon the server to limit results by a function comparing the distribution of their response times

JSON response data dominated the requests, being the only format available across all three server types. XML's geographic subset, GML, is only routinely served by OGC endpoints. Images were not requested as often, there being fewer server-side filtering functions available.



Users may request the latest pie chart from the LandgateAPITest web app at this [link](#).

Figure 10 - Percentage of Vectors by response data type

The box and whisker chart shows the response time distribution for XML responses is tighter and higher overall than the similar JSON and image request response time distributions. All three have a significant number of outliers in their upper ranges, showing clearly skewed distributions with most requests achieved in short time frames.

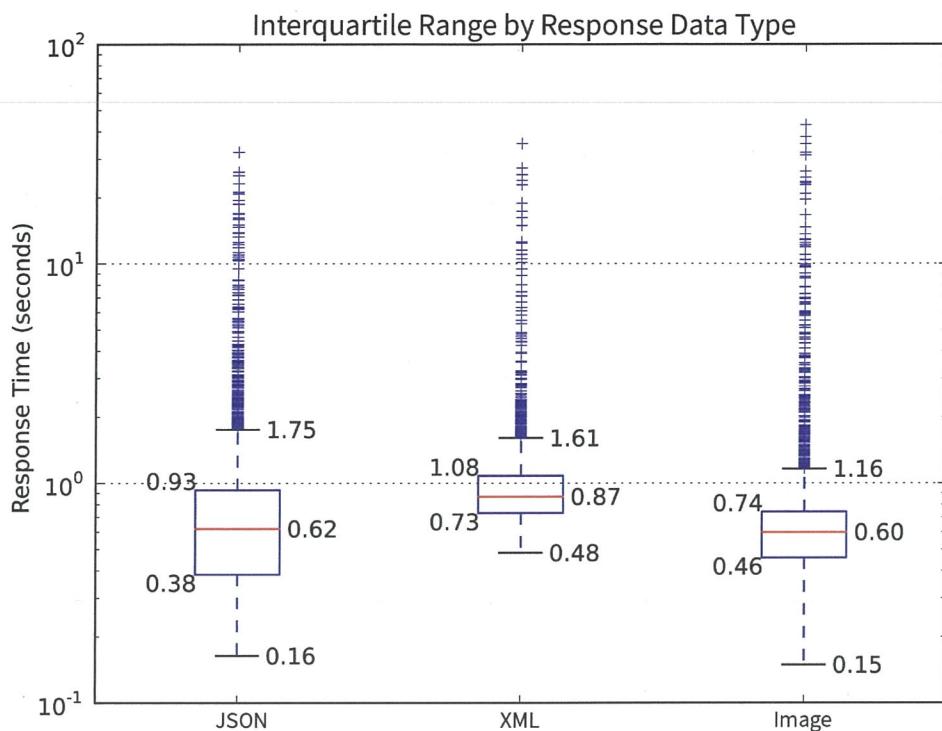


Figure 11 - Response time distribution by response data type

The Esri and OGC portion of the test campaign in March 2016 was more vigorous than the earlier GME part in December 2015. This pie chart shows how clearly the Esri and OGC tests overwhelm the fewer GME ones.

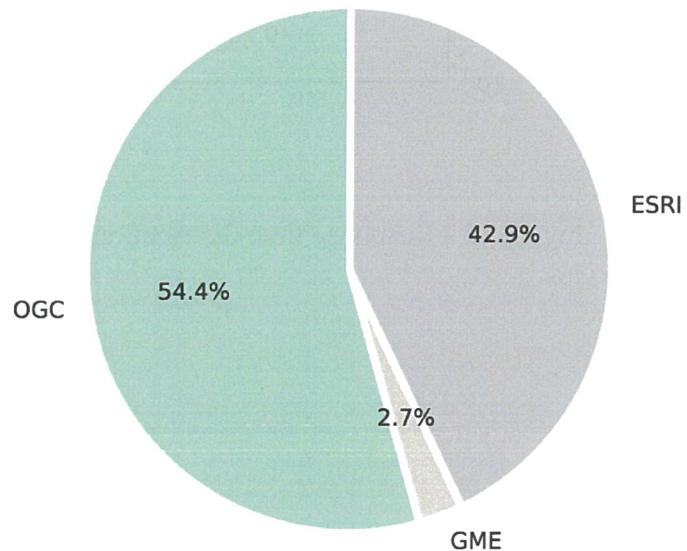


Figure 12 - Percentage of Vector objects by server type

The box and whiskers chart appears to show a clear performance win for the Esri servers over the OGC servers, having a much lower median and interquartile range. There is a significant consideration here that Esri servers do not supply heavier

payload XML/GML responses where OGC ones do. As the response data type and response size charts show, on average larger responses have slower response times.

The GME tests fill a broader interquartile range and have fewer outliers. A larger sample set of these requests ~~may~~ could have increased our confidence in this result.

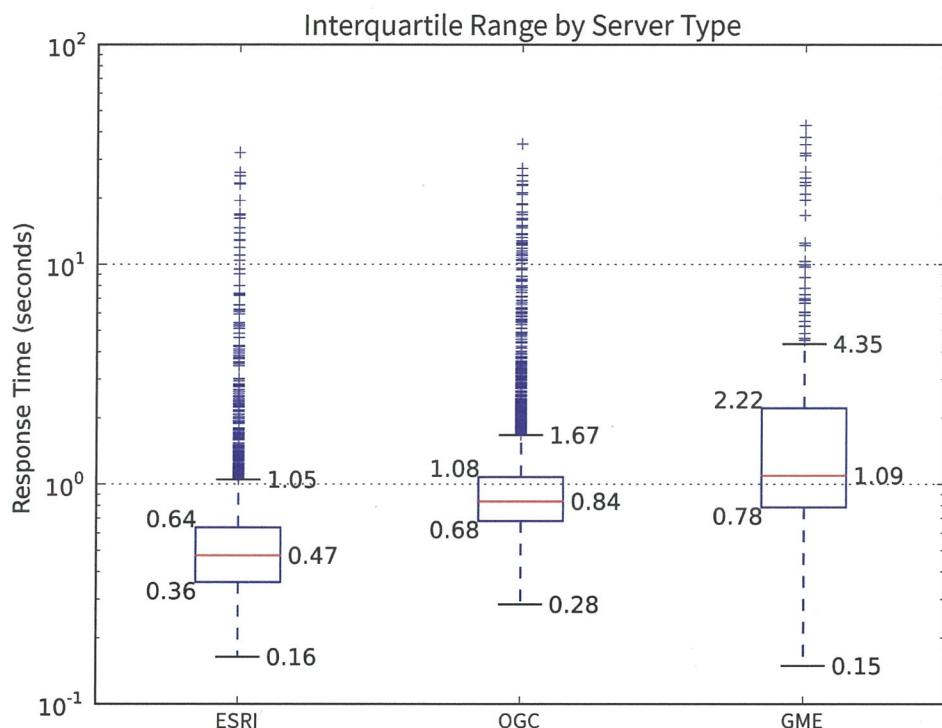


Figure 13 - Response time distribution by server type

This pie chart shows tests were almost evenly split between the two HTTP methods favoured by spatial servers; GET and POST. The greater proportion of GET requests are partly due to the lack of POST requests created for the GME server and the map tile requests mostly being GET's with key value coding or straight restful endpoints.

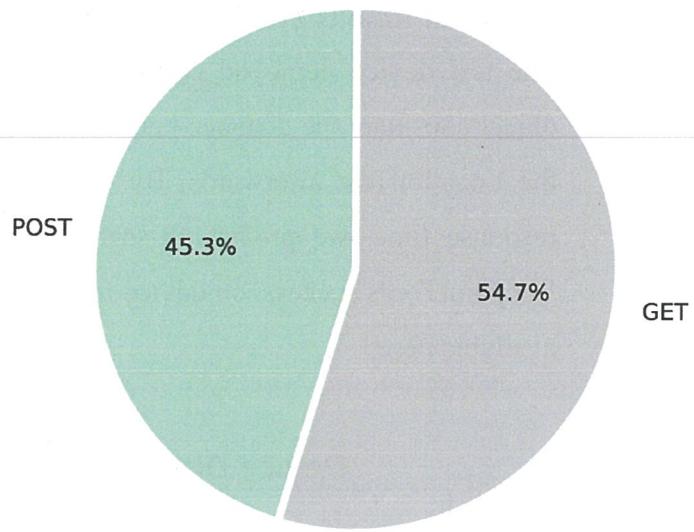


Figure 14 - Percentage of Vectors by HTTP Method (GET and POST)

There was no distinct difference in response time between the two methods. The medians and boundaries of interquartile ranges are similar enough that differences could be rounding errors. The latest graph is available here.

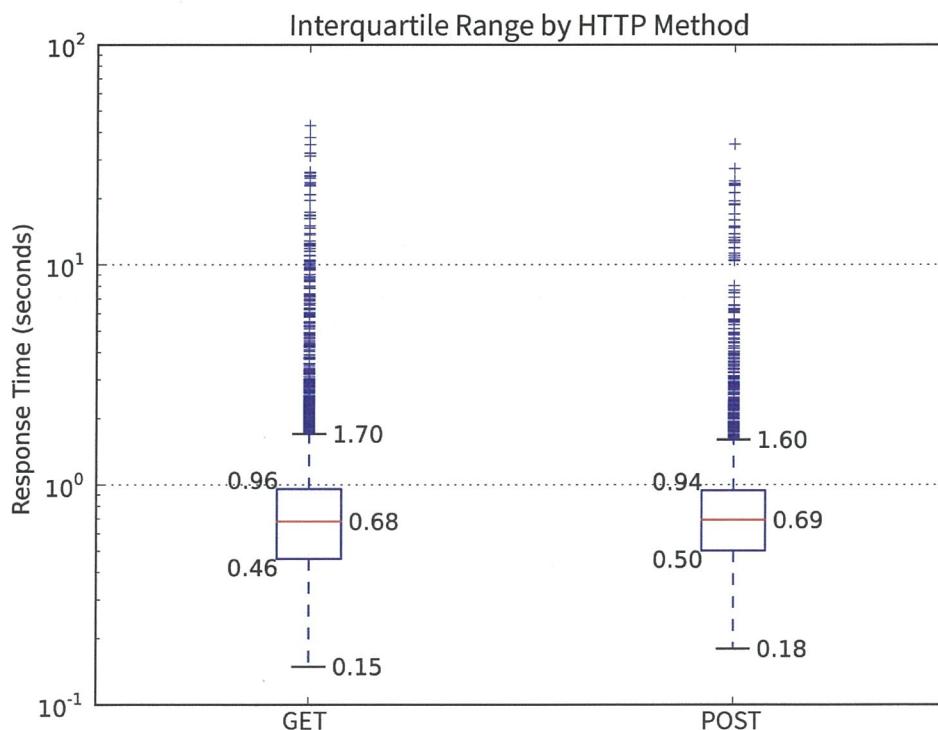


Figure 15 - HTTP Method (GET and POST) response time distributions in box plot

#### 4.5 Test Results by Distance Device Travelled

The test device deployed could determine its location through GPS. The Vector object considers the distance between the LocationTest prior to an EndpointTest and the LocationTest afterwards. By comparing each Vector's distance property to its response time, we produce a scatter plot. Then we categorise the points by the EndpointTest's success, on device failure or reference check failure. The live graph is available [here](#).

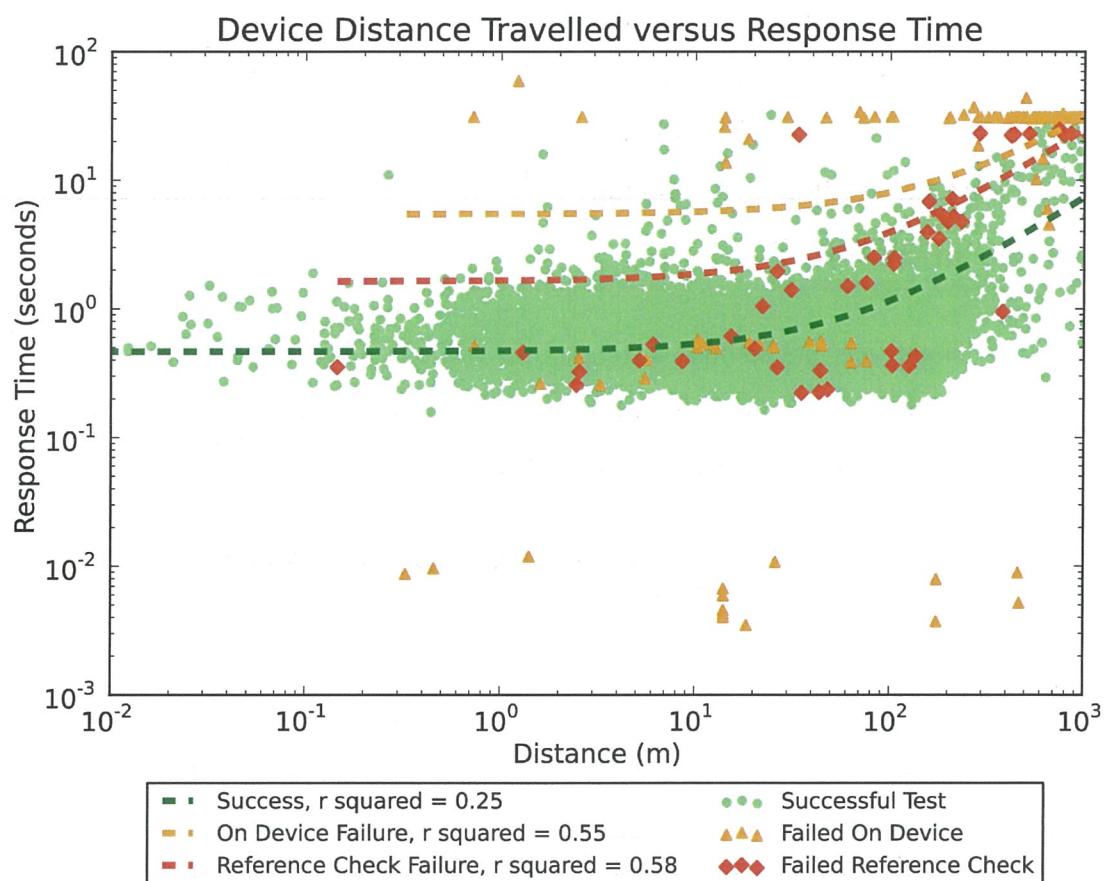


Figure 16 - A scatterplot of distance device travelled (metres) versus response time (seconds) for each Vector object

The green successful tests show a loose trend of increasing response time in line with increasing distance travelled.

The orange "failed on device" category exhibit three distinct horizontal bands of response times. The bottommost band are the shortest response times, notably so as they are at most 1/100th of a second. These are the cases where the device had no

mobile network connection at all and aborted the request immediately without even an attempt to send it to the server.

The middle band of "on device failures" are those with similar response times to many successful requests. These are the tests cancelled before completion. They had a successful link to the server but the test was interrupted by an incoming phone call or the app was otherwise switched to the background. In line with the application's design goals these tests were aborted, their response time recorded and marked as on device failures.

The uppermost band forms a clear line around 30 seconds in response time. This is the standard time-out length for a web service request on an iOS device. Requests without a response are aborted by the system. Interestingly, the majority of these failures occurred when the device travelled more than 100m.

The tests which failed their reference check appear to follow a similar trend to the successful tests. This is due to the fact that their web service requests were able to complete successfully from the device's point of view.

The scatterplot shows enough noise to produce R-squared values that are less than ideal. Each value of distance produces a range of response times due to several factors, most notably the response payload size varies by request type. Distance values are not perfect either as we must give consideration to the GPS receiver's desired and possible accuracy.